

Everyone got the tool. Not everyone got the setup instructions.

Workbrain enables participants of any skill level to write the instructions that make any AI tool work like it already knows them — **their words, their standards, their way of doing the work** — and they prove it against their own before-and-after. They keep the file and the knowledge of how to leverage it towards their goals and objectives.

4-Week Training Program

90 Minutes Per Week

We Facilitate

You Host

Free For Participants

 **What It Costs**

Hope this email finds you —


You filled that in before I finished the sentence. **So does the tool everyone in your building is already using** — for every word it writes.

Hope this email finds you well |


well



in good health




rested after the strike



The AI tools we all used today are a lot like autocomplete, at enormous scale, **answering our prompts with the most common response.**

The most common version of the thought is becoming **the noise**.

The most likely word is the most common one — that isn't a flaw, it's the definition. So every draft that comes back starts life as the most common version of the thought. **Multiply that by everyone, every day.**



It is already getting hard to tell where the tool ended and the person began — in an inbox, in an application, in a proposal. **That line is worth being able to draw.**

The average has an **accent**.

"Most likely" is measured over text nobody collected evenly — so the middle belongs to whoever wrote the most of it. Watch what that does to a sentence somebody meant:

WHAT THEY SAID

"We torque it to spec, then check it again cold."

"I've run this line eleven years — I hear it before the gauge does."

"The crew calls it the ghost shift, and they plan around it."

WHAT COMES BACK

→ *We follow standard quality procedures.*

→ *Experienced operations professional.*

→ *The overnight production period.*

Nothing was flagged. Nobody was asked. **The words that carried the expertise are the ones it edits hardest** — and what comes back could have been written about anyone, at any company.



A **Workbrain file** makes their words the starting point — so the person whose knowledge is least like everyone else's gets the same leverage as the person nearest the middle.

SOURCES Hofmann et al., Nature 2024 — models judged people more harshly for writing in a dialect, and debiasing training mostly hid it. Santurkar et al., ICML 2023 — model defaults match some groups' views far more closely than others'.

And it **reads** them the same way.

Interpretation is prediction too. Handed a request from the ops lead who has been with you fifteen years, it hears the **average asker's** meaning — not hers:

WHAT THEY ASKED

"Write up what the crew flagged on the Johnson job."

"Draft the note about the ghost shift."

"Keep it tight for the shop floor."

WHAT IT HEARD

→ *A thank-you letter to a customer named Johnson.*

→ *A memo on overnight staffing costs.*

→ *Bulleted, formal, and missing the one caveat that matters.*


It does not ask. It answers the question it assumed — **confidently, and in the same breath**. The shorthand it drops is the part a person would have known to keep.

SOURCE Sclar et al., ICLR 2024 — rewording the same question moved measured accuracy by up to 76 points.

05 · THEY'VE SEEN THIS MOVIE

A profile about you, written by **someone else.**

Twenty years of web cookies: a dossier nobody authored, kept for someone else's purposes, invisible to the person it described. **We ended up regulating it.** Your customers remember how that felt. So does your team.



Every one of them already has one.



COMMON

On by default on all five major services surveyed. Nobody opted in.



ASSUMED

Inferred, not written — and it rewrites itself over time. OpenAI calls that process "dreaming."



LOCKED

At Microsoft: **exempt from retention policy, not audit-logged, not restrictable** — their own documentation.



RENTED

At Perplexity Enterprise those memories **belong to the organization** — access ends when the person leaves.

Company work is drafted through profiles nobody here authored, audited, or can read in full — and **no vendor will say which one wins** when it disagrees with your own instructions.

Three questions. No cost, no sign-up.

ACCURACY

"What do you remember about me?"

Look for a job title, a project, or a tool you have already moved on from.

Confirm it by opening the memory page itself — if it is still listed there, it is still shaping every answer.

Field guide: last slide.

VOICE

"Describe how I communicate."

Look for an adjective you would never use about yourself.

Confirm it by holding the description next to something you actually wrote last week.

Field guide: last slide.

PRIORITIES

"What matters most in my work?"

Look for what is missing — the thing you would have said first.

Confirm it by asking it to rank three of your real responsibilities, then check the order against yours.

Field guide: last slide.

Send them to your team, or try them yourself tonight. **Whatever happens on your screen is happening on theirs too** — every day, on the accounts drafting your client work.

The effort doesn't change. The **ceiling** does.

Nobody accepts that memory on principle — it just costs nothing and comes back well-composed. And **we all paste straight from the tool more often than we'd admit.** Keep it.

DAILY EFFORT

Memory alone

none

After four sessions

about the same

WHAT THE TOOL CAN BE TRUSTED WITH

Memory alone

drafts, in the house style

After four sessions

drafts in their voice · repeatable methods · actions worth automating

09 · THE MEASURED GROUP EFFECT

Individually faster. Collectively **identical**.

Each person's output improves, and **the group's diversity shrinks**. Ten people on one team produce ten versions of the same thing — and none of them chose that.




SOURCES Doshi & Hauser, Science Advances 2024 · Padmakumar & He, ICLR 2024 · Shumailov et al., Nature 2024. A measured trajectory, still early — we say so out loud.

10 · WHAT IT COSTS THEM


The day their work stops sounding like **them**.

For the person: the contribution that used to be recognizably hers stops getting noticed. For you: in a stack where every proposal is polished by the same tool, **sameness is how you lose a deal you used to win.**



EVERY DRAFT, POINTING AT EVERY OTHER DRAFT

The difference has to come from the **people**.



Your competitor's proposal was written by the same model as yours. So was the pitch from the shop across town. **Has anyone given your team a way to put themselves back into the work?**

Sit in their seat for a second.

Someone on your team opens the same tool everyone else has. Here is the whole menu available to them today:

Avoid it

Stay away from it. Reasonable — they have heard it is unreliable, or that IT has not blessed it. Meanwhile the people they compete with are twice as fast, and nobody tells them that is what is happening.

Wing it

Use it for everything, exactly as it comes. Also reasonable — it is right there and it is fast. Their work now sounds like everyone's work, and the expertise you hired them for never makes it onto the page.

Those are the two choices they actually have. Neither is a plan, and nobody has offered them a third. **Four sessions is the third.**

13 · THE THIRD OPTION

Teach them to **author** their own context.

Each person writes a short, plain-text file about their own work — their names for things, their judgment, the standards they hold — and hands it to whatever tool they already use. **Nothing routes through us. No account. Nothing transmitted.**

They leave with a simple file. A plain text file. **A configured, plain text file.** We have all gone through the setup process on a new device or a new tool. Workbrain is the setup for effective AI usage.

A skill that belongs to them.

Their files

Plain markdown they own — who they are, how they work, what they repeat. On their device, not ours.

A method

One job they do every week, written down well enough that a tool — or a new coworker — can follow it.

Proof

Their own before-and-after, run with their own eyes, on their own work.

Institutional memory

The way your best person actually does the job, written down — so when they are out, or gone, **the method does not leave with them.**

It works in a tool that doesn't exist yet, in a model nobody has released. The files belong to each person — and **the written methods stay with the team.**

Numbers, not testimonials.

You can say what "good" means

Four plain checks anyone can run in five minutes, published on a public page. Your most skeptical engineer can audit the claim without asking us anything.

You can say how we know

Every participant does the same kind of task before any teaching and again at the end, and keeps both. The evidence is on their own work — real tasks from your shop, not a survey about how the training felt.

You can say what we won't claim

At 90 and 180 days we report back, always with the response rate attached. We do not claim employment outcomes — people tell us what happened, and we pass it on as reported.

All of it is public, including the six claims this program rests on — written down so anyone who thinks we are wrong can say exactly where: myworkbrain.org/good

One number — and more sessions barely move it.

PROGRAM BASE — PAID ONCE, WHATEVER YOU RUN

Curriculum, facilitator preparation, setup	\$1,000
Loaner device pool, shared across every session	\$900
Insurance and administration	\$500
Follow-up reporting at 90 and 180 days	\$500

EACH SERIES — 4 SESSIONS · UP TO 10 SEATS

Facilitation — on site, morning, afternoon or evening	\$1,600
Printed materials and certificates	\$300

Your first series, all in **\$4,800**

Every additional series — any day, any time slot **+\$1,900**

RUN MORE, PAY LESS PER PERSON

\$480

per seat
one team of 10

\$335

per seat
two teams · 20 people

\$287

per seat
three teams · 30 people

WHAT IS NOT ON THIS INVOICE

- No cost to anyone in the room, and no charge for a person who has to miss a session.
- No software license, seat fee or subscription — before, during, or after.
- No second invoice to launch: the base covers setup, devices and reporting for every series you run.
- Nothing lapses when we leave. What they build keeps working, free, permanently.

From you: a room with wifi, a repeating slot that suits the shift pattern, and the people who should be in it. We do not sell to your staff. Insurance carried, facilitators background-checked, documentation on request.

Take it, or keep the three questions.

If this is something you want for your team: ask for a series at myworkbrain.org/services and the outline, the agreement and available dates come back within two business days. If it isn't: **keep the three questions, run them at your next staff meeting, and we are square.** There is no follow-up sequence.

APPENDIX A · PUSH BACK ON US

"The homogenization research is early."

It is, and we say so on the slide. What isn't early: memory on by default, profiles nobody authored, and the before-and-after each person measures on their own work. The case never rests its weight on the scary slide.

"This sounds anti-AI. We just rolled it out company-wide."

Four 90-minute sessions of hands-on building with the tool each person already uses. The foundations hour exists so the adoption is informed. We teach use — the caution is the seatbelt, not an argument against the car.

"My team needs output now, not voice lessons."

Everyone's output is polished now, and buyers pattern-match the house style — in proposals, in pitches, in hiring. When polish is table stakes, what survives is whose judgment shows. That is a revenue argument, not a taste argument.

APPENDIX B · THE HONEST CONCESSIONS

"Isn't this prompt training with extra steps?"

Prompt phrasing doesn't transfer between tools — the research is clear. An owned file does. The room demonstrates the difference to itself, in five minutes, on their own work.

"You teach people not to sound like the machine... using the machine."

The same way a writing class uses a word processor. The tool drafts; the person and their file decide. Improve what they do — don't replace it.

What we concede without a fight

Average is fine for boilerplate; the file exists for the work that carries someone's name. And a context file steers voice — it does not guarantee it. The promise is authorship plus evidence, never magic.

Where the memory lives.

Claude

Settings → memory. Viewable by topic; import and export are documented.

ChatGPT

Settings → Personalization → Memory. A summary page; background synthesis keeps updating it.

Gemini

Saved Info. Note that its memory and its Gems are mutually exclusive surfaces.

Microsoft 365 Copilot

Settings → personalization. Stored in the person's Exchange mailbox; on by default.

Perplexity

Settings → memory. In Enterprise, those memories belong to the organization.

Paths current as of September 2026 — vendors move them; search "<service> memory settings" if one has drifted. And remember: the page shows only the part they show you.