

Everyone gave them the tool. Nobody taught them to **own it.**

Ten minutes on a skill you can give your team — one that works in any tool, survives whichever app is popular this year, and leaves the method behind when the person moves on.

\$4,800 funds a first series — four sessions, up to 10 of your team, mornings or evenings · the price is on this slide because there is no follow-up sequence

01 · WHAT THEY'RE ACTUALLY HOLDING

Hope this email finds you _____

You filled that in before I finished the sentence. **So does the tool everyone in your building is already using** — for every word it writes.

The AI tools we all used today are a lot like autocomplete, at enormous scale, answering our prompts with the most common response.

A simplification an expert accepts: predicting the most likely next word is the literal training objective.

The most common version of the thought is becoming **the noise.**

The most likely word is the most common one — that isn't a flaw, it's the definition. So every draft that comes back starts life as the most common version of the thought.

Multiply that by everyone, every day.

It is already getting hard to tell where the tool ended and the person began — in an inbox, in a proposal, in a quarterly update. **That line is worth being able to draw.**

The average has an **accent**.

"Most likely" is measured over text that was never collected evenly from everyone — so the middle belongs to whoever wrote the most of it. Watch what that does to a sentence somebody meant:

WHAT THEY SAID

"We torque it to spec, then check it again cold."

"I've run this line eleven years — I hear it before the gauge does."

"The crew calls it the ghost shift, and they plan around it."

WHAT COMES BACK

→ *We follow standard quality procedures.*

→ *Experienced operations professional.*

→ *The overnight production period.*

Nothing was flagged. Nobody was asked. **The words that carried the expertise are the ones it edits hardest** — and what comes back could have been written about anyone, at any company.

An authored file is how the person whose knowledge is least like everyone else's gets the same leverage as the person nearest the middle.

Measured, not asserted: covert prejudice from dialect alone, surviving the training meant to remove the overt kind (Hofmann et al., Nature 2024); model defaults aligning far more closely with some demographic groups than others (Santurkar et al., ICML 2023). A file does not fix a model — it puts the person's own words in front of it.

And it **reads** them the same way.

Interpretation is prediction too. Handed a request from the ops lead who has been with you fifteen years, it hears the most likely meaning — the **average asker's** meaning.

What it misreads is exactly what you pay her for: her judgment, her shorthand, the fifteen years of things she knows that nobody else does. **And it never says so.** It answers a question nobody asked, confidently.

Rewording alone swings measured accuracy up to 76 points (Sclar et al., ICLR 2024). Same engine, pointed the other way: the accent on the previous slide is also an accent it reads with.

05 · THEY'VE SEEN THIS MOVIE

A profile about you, written by **someone else.**

Twenty years of web cookies: a dossier nobody authored, kept for someone else's purposes, invisible to the person it described. **We ended up regulating it.** Your customers remember how that felt. So does your team.

Every one of them already has one.



COMMON

On by default on all five major services surveyed. Nobody opted in.



ASSUMED

Inferred, not written — and it rewrites itself over time. OpenAI calls that process "dreaming."



LOCKED

At Microsoft: **exempt from retention policy, not audit-logged, not restrictable** — their own documentation.



RENTED

At Perplexity Enterprise those memories **belong to the organization** — access ends when the person leaves.

Receipts cited in "Durable Context" (Young, 2026) — myworkbrain.org. Company work is being drafted through profiles nobody at your company authored, audited, or can read in full — and no vendor will tell you which one wins when it disagrees with your own instructions.

Three questions. No cost, no sign-up.

ACCURACY

"What do you remember about me?"

Look for a job title, a project, or a tool you have already moved on from.

Confirm it by opening the memory page itself — if it is still listed there, it is still shaping every answer.

Field guide: last slide.

VOICE

"Describe how I communicate."

Look for an adjective you would never use about yourself.

Confirm it by holding the description next to something you actually wrote last week.

Field guide: last slide.

PRIORITIES

"What matters most in my work?"

Look for what is missing — the thing you would have said first.

Confirm it by asking it to rank three of your real responsibilities, then check the order against yours.

Field guide: last slide.

Send them to your team, or try them yourself tonight. Whatever happens on your screen is happening on theirs too, every day, quietly — on the accounts drafting your client work.

The effort doesn't change. The **ceiling** does.

Nobody accepts that memory on principle — it just costs nothing and comes back well-composed. And **we all paste straight from the tool more often than we'd admit.** Keep it.

DAILY EFFORT

Memory alone

none

After four sessions

about the same

WHAT THE TOOL CAN BE TRUSTED WITH

Memory alone

drafts, in the house style

After four sessions

drafts in their voice · repeatable methods · actions worth automating

One hour of writing, once, barely moves tomorrow's effort. It moves what the tool is allowed to do — and written methods are the prerequisite for automating anything responsibly.

09 · THE MEASURED GROUP EFFECT

Individually faster. Collectively **identical**.

The research finding on AI assistance: each person's output improves, and **the group's diversity shrinks**. Ten people on one team produce ten versions of the same thing — and none of them chose that.

Doshi & Hauser, Science Advances 2024 · Padmakumar & He, ICLR 2024 · Shumailov et al., Nature 2024. A measured trajectory, still early — we say so out loud.

10 · WHAT IT COSTS THEM

The day their work stops sounding like **them**.

For the person: the contribution that used to be recognizably hers stops getting noticed. For you: in a stack where every proposal is polished by the same tool, **sameness is how you lose a deal you used to win.** The discount lands on your brand, never on the software.

11 • ONE QUESTION TO KEEP

The difference has to come from the **people.**

Your competitor's proposal was written by the same model as yours. So was the pitch from the shop across town. **Has anyone given your team a way to put themselves back into the work?**

You have two choices to make right now.

Avoidance

Hold the line until it settles down, or until Legal signs off. Perfectly reasonable — and they are using it anyway, on personal accounts, on their phones, with nobody to ask. You get none of the oversight and none of the learning.

Acceptance

Let everyone get on with it. Also reasonable — and unsupported adoption walks the whole room toward the same house style, spending the very thing you hired each of them for.

Both are plausible. Both leave everyone to work it out alone. **And the organizations that look miles ahead of you chose neither** — they taught their people a third thing, and your team is four sessions from the same unlock. They will credit it to you.

13 · THE THIRD OPTION

Teach them to **author** their own context.

Each person writes a short, plain-text file about their own work — their names for things, their judgment, the standards they hold — and hands it to whatever tool they already use. **Nothing routes through us. No account. Nothing transmitted.**

They leave with a simple file. A plain text file. **A configured, plain text file.** We have all gone through the setup process on a new device or a new tool. Workbrain is the setup for effective AI usage.

A skill that belongs to them.

Their files

Plain markdown they own — who they are, how they work, what they repeat. On their device, not ours.

A method

One job they do every week, written down well enough that a tool — or a new coworker — can follow it.

Proof

Their own before-and-after, run with their own eyes, on their own work.

Institutional memory

The way your best person actually does the job, written down — so when they are out, or gone, **the method does not leave with them.**

It works in a tool that doesn't exist yet, in a model nobody has released. The files belong to each person — and **the written methods stay with the team.**

Numbers, not testimonials.

You can say what "good" means

Four plain checks anyone can run in five minutes, published on a public page. Your most skeptical engineer can audit the claim without asking us anything.

You can say how we know

Every participant does the same kind of task before any teaching and again at the end, and keeps both. The evidence is on their own work — real tasks from your shop, not a survey about how the training felt.

You can say what we won't claim

At 90 and 180 days we report back, always with the response rate attached. We do not claim employment outcomes — people tell us what happened, and we pass it on as reported.

All of it is public, including the six claims this program rests on — written down so anyone who thinks we are wrong can say exactly where: myworkbrain.org/good

One number — and more sessions barely move it.

PROGRAM BASE — PAID ONCE, WHATEVER YOU RUN

Curriculum, facilitator preparation, setup	\$1,000
Loaner device pool, shared across every session	\$900
Insurance and administration	\$500
Follow-up reporting at 90 and 180 days	\$500

EACH SERIES — 4 SESSIONS · UP TO 10 SEATS

Facilitation — on site, morning, afternoon or evening	\$1,600
Printed materials and certificates	\$300

Your first series, all in **\$4,800**

Every additional series — any day, any time slot **+\$1,900**

RUN MORE, PAY LESS PER PERSON

\$480

per seat
one team of 10

\$335

per seat
two teams · 20 people

\$287

per seat
three teams · 30 people

WHAT IS NOT ON THIS INVOICE

- No cost to anyone in the room, and no charge for a person who has to miss a session.
- No software license, seat fee or subscription — before, during, or after.
- No second invoice to launch: the base covers setup, devices and reporting for every series you run.
- Nothing lapses when we leave. What they build keeps working, free, permanently.

From you: a room with wifi, a repeating slot that suits the shift pattern, and the people who should be in it. We do not sell to your staff. Insurance carried, facilitators background-checked, documentation on request.

Take it, or keep the three questions.

If this is something you want for your team: ask for a series at myworkbrain.org/services and the outline, the agreement and available dates come back within two business days. If it isn't: **keep the three questions, run them at your next staff meeting, and we are square.** There is no follow-up sequence.

APPENDIX A · PUSH BACK ON US

"The homogenization research is early."

It is, and we say so on the slide. What isn't early: memory on by default, profiles nobody authored, and the before-and-after each person measures on their own work. The case never rests its weight on the scary slide.

"This sounds anti-AI. We just rolled it out company-wide."

Four 90-minute sessions of hands-on building with the tool each person already uses. The foundations hour exists so the adoption is informed. We teach use — the caution is the seatbelt, not an argument against the car.

"My team needs output now, not voice lessons."

Everyone's output is polished now, and buyers pattern-match the house style — in proposals, in pitches, in hiring. When polish is table stakes, what survives is whose judgment shows. That is a revenue argument, not a taste argument.

APPENDIX B · THE HONEST CONCESSIONS

"Isn't this prompt training with extra steps?"

Prompt phrasing doesn't transfer between tools — the research is clear. An owned file does. The room demonstrates the difference to itself, in five minutes, on their own work.

"You teach people not to sound like the machine... using the machine."

The same way a writing class uses a word processor. The tool drafts; the person and their file decide. Improve what they do — don't replace it.

What we concede without a fight

Average is fine for boilerplate; the file exists for the work that carries someone's name. And a context file steers voice — it does not guarantee it. The promise is authorship plus evidence, never magic.

Where the memory lives.

Claude

Settings → memory. Viewable by topic; import and export are documented.

ChatGPT

Settings → Personalization → Memory. A summary page; background synthesis keeps updating it.

Gemini

Saved Info. Note that its memory and its Gems are mutually exclusive surfaces.

Microsoft 365 Copilot

Settings → personalization. Stored in the person's Exchange mailbox; on by default.

Perplexity

Settings → memory. In Enterprise, those memories belong to the organization.

Paths current as of September 2026 — vendors move them; search "<service> memory settings" if one has drifted. And remember: the page shows only the part they show you.