

Everyone gave them the tool. Nobody taught them to **own it.**

Ten minutes on a skill you can put in the hands of the people who count on you — one that stays theirs, works in any assistant, and outlives whichever app is popular this year.

\$4,800 funds one cohort of 16–20 — your staff, your members, or a mix · they never pay a cent · the price is on this slide because there is no follow-up sequence

01 · WHAT THEY'RE ACTUALLY HOLDING

Hope this email finds you _____

You filled that in before I finished the sentence. **So does the assistant everyone in your building is already using** — for every word it writes. It is autocomplete at enormous scale: it continues with what most people would say.

A simplification an expert accepts: predicting the most likely next word is the literal training objective.

The most common version of the thought is becoming **the noise.**

The most likely word is the most common one — that isn't a flaw, it's the definition. So every draft that comes back starts life as the most common version of the thought.

Multiply that by everyone, every day.

It is already getting hard to tell where the assistant ended and the person began — in an inbox, in an application, in a proposal. **That line is worth being able to draw.**

03 · THE HALF NOBODY MENTIONS

And it **reads** them the same way.

Interpretation is prediction too. Handed a request from the coordinator who runs your Tuesday programs, it hears the most likely meaning — the **average asker's** meaning. What it misreads is exactly what you value in her: her judgment, her shorthand, the fifteen years of things she knows that most people don't. And it never says so. It answers a question nobody asked, confidently.

Rewording alone swings measured accuracy up to 76 points (Sclar et al., ICLR 2024). Models misread — and covertly penalize — non-standard dialects (Hofmann et al., Nature 2024). It does not read everyone equally well — and the people it reads worst are often the ones you are working hardest to help.

04 · THEY'VE SEEN THIS MOVIE

A profile about you, written by **someone else.**

Twenty years of web cookies: a dossier nobody authored, kept for someone else's purposes, invisible to the person it described. **We ended up regulating it.** Everybody remembers how that felt.

Every one of them already has one.

- AI memory is **on by default** on all five major services surveyed.
- **Inferred, not written by them** — and it rewrites itself over time (OpenAI calls the process "dreaming").
- At Microsoft: **exempt from retention policy, not audit-logged, not restrictable** — their own documentation.
- At Perplexity Enterprise: those memories **belong to the organization**, and access ends when the person leaves.

Receipts cited in "Durable Context" (Young, 2026) — myworkbrain.org. Nobody wrote the profile their assistant reads them through — not the person it describes, not you. Most have never laid eyes on it.

Three questions. No cost, no sign-up.

ACCURACY

"What do you remember about me?"

Look for a job title, a project, or a tool you have already moved on from.

Confirm it by opening the memory page itself — if it is still listed there, it is still shaping every answer.

Field guide: last slide.

VOICE

"Describe how I communicate."

Look for an adjective you would never use about yourself.

Confirm it by holding the description next to something you actually wrote last week.

Field guide: last slide.

PRIORITIES

"What matters most in my work?"

Look for what is missing — the thing you would have said first.

Confirm it by asking it to rank three of your real responsibilities, then check the order against yours.

Field guide: last slide.

Send them to your staff, post them for your members, or try them yourself tonight. Whatever happens on your screen is happening on theirs too, every day, quietly.

The effort doesn't change. The **ceiling** does.

Nobody accepts that memory on principle — it just costs nothing and comes back well-composed. And **we all paste straight from the assistant more often than we'd admit.** Keep it.

DAILY EFFORT

Memory alone

none

After four evenings

about the same

WHAT THE TOOL CAN BE TRUSTED WITH

Memory alone

drafts, in the house style

After four evenings

drafts in their voice · repeatable methods · actions worth automating

One hour of writing, once, barely moves tomorrow's effort. It moves what the tool is allowed to do — and written methods are the prerequisite for automating anything responsibly.

08 · THE MEASURED GROUP EFFECT

Individually faster. Collectively identical.

The research finding on AI assistance: each person's output improves, and **the group's diversity shrinks**. Ten people produce ten versions of the same thing — and none of them chose that.

Doshi & Hauser, Science Advances 2024 · Padmakumar & He, ICLR 2024 · Shumailov et al., Nature 2024. A measured trajectory, still early — we say so out loud.

09 · WHAT IT COSTS THEM

The day their work stops sounding like **them**.

For someone on your staff: the contribution that used to be recognizably hers stops getting noticed. For anyone you serve who is job-hunting, in a pile where every application is polished by the same tool, **sameness is the rejection**. And the discount lands on your organization's name, never on the software.

10 • ONE QUESTION TO KEEP

The difference has to come from the **people.**

Your competitor's proposal was written by the same model as yours. So was the grant application from the organization across town. **Has anyone given the folks doing that work a way to put themselves back into it?**

You have two choices to make right now.

Avoidance

Hold the line until it settles down. Perfectly reasonable — and they are using it anyway, on personal accounts, at home, with nobody to ask. You get none of the oversight and none of the learning.

Acceptance

Let everyone get on with it. Also reasonable — and unsupported adoption walks the whole room toward the same house style, spending the very thing that made each person worth hiring.

Both are plausible. Both leave everyone to work it out alone. **And the organizations that look miles ahead of you chose neither** — they taught their people a third thing, and the folks you serve are four evenings from the same unlock. They will credit it to you.

12 · THE THIRD OPTION

Teach them to **author** their own context.

Each person writes a short, plain-text file about their own work — their names for things, their judgment, the standards they hold — and hands it to whatever assistant they already use. **Nothing routes through us. No account. Nothing transmitted.** The tool finally reads and writes from who they actually are.

A skill that belongs to them.

Their files

Plain markdown they own — who they are, how they work, what they repeat. On their device, not ours.

A method

One job they do every week, written down well enough that a tool — or a new coworker — can follow it.

Proof

Their own before-and-after, run with their own eyes, on their own work.

A credential

A certificate at session four, and **Work Brain Certified** later for demonstrated use — verifiable, and it goes on a profile.

It works in a tool that doesn't exist yet, at their next job, in a model nobody has released. **Nothing here expires when the program ends.**

Numbers, not testimonials.

You can say what "good" means

It is four plain checks anyone can run in five minutes, published on a public page.

Nobody has to take our word for what a good result looks like — including you.

You can say how we know

Every participant does the same kind of task before any teaching and again at the end, and keeps both. The evidence is theirs, on their own work — not a survey about how the class felt.

You can say what we won't claim

At 90 and 180 days we report back, always with the response rate attached. We do not claim employment outcomes — people tell us what happened, and we pass it on as reported.

All of it is public, including the six claims this program rests on — written down so anyone who thinks we are wrong can say exactly where: myworkbrain.org/good

One number, and it is the whole number.

Facilitation — four evenings and a board night	\$2,600
Curriculum, printed materials, certificates	\$400
Loaner devices for anyone without one	\$900
Insurance and administration	\$500
Follow-up reporting at 90 and 180 days	\$400
One cohort · 16–20 seats	\$4,800

WHAT IS NOT ON THIS INVOICE

- No cost to the people in the room — not a registration fee, not a book, not a device.
- No software license, no seat fee, no subscription — before, during, or after.
- No second invoice to launch: materials, certificates, devices and reporting are in the number above.
- Nothing that lapses when we leave. What they build keeps working for free, permanently.
- From you: a room with wifi, the same evening each week, and the folks who should be in it.

We do not sell to participants, during or after. Liability insurance carried, facilitators background-checked, documentation on request. Taster and train-your-own-facilitator formats are in the packet.

Take it, or keep the three questions.

If this is something you want for the people you serve: request a cohort at myworkbrain.org/civic and the packet, the agreement and available dates come back within two business days. If it isn't: **keep slide six, run it with your staff or your members, and we are square.** There is no follow-up sequence.

APPENDIX A · PUSH BACK ON US

"The homogenization research is early."

It is, and we say so on the slide. What isn't early: memory on by default, profiles nobody authored, and the before-and-after each person measures on their own work. The case never rests its weight on the scary slide.

"This sounds anti-AI. We're trying to drive adoption."

Four evenings of hands-on building with each person's own assistant. The foundations hour exists so the adoption is informed. We teach use — the caution is the seatbelt, not an argument against the car.

"The folks I serve need results now, not voice lessons."

Everyone's output is polished now, and readers pattern-match the house style — in hiring, in sales, in grant review. When polish is table stakes, what survives is whose judgment shows. That is an output argument, not a taste argument.

APPENDIX B · THE HONEST CONCESSIONS

"Isn't this prompt training with extra steps?"

Prompt phrasing doesn't transfer between assistants — the research is clear. An owned file does. The room demonstrates the difference to itself, in five minutes, on their own work.

"You teach people not to sound like the machine... using the machine."

The same way a writing class uses a word processor. The tool drafts; the person and their file decide. Improve what they do — don't replace it.

What we concede without a fight

Average is fine for boilerplate; the file exists for the work that carries someone's name. And a context file steers voice — it does not guarantee it. The promise is authorship plus evidence, never magic.

Where the memory lives.

Claude

Settings → memory. Viewable by topic; import and export are documented.

ChatGPT

Settings → Personalization → Memory. A summary page; background synthesis keeps updating it.

Gemini

Saved Info. Note that its memory and its Gems are mutually exclusive surfaces.

Microsoft 365 Copilot

Settings → personalization. Stored in the person's Exchange mailbox; on by default.

Perplexity

Settings → memory. In Enterprise, those memories belong to the organization.

Paths current as of September 2026 — vendors move them; search "<service> memory settings" if one has drifted. And remember: the page shows only the part they show you.