

Everyone gave them the tool. Nobody taught them to **own it.**

Ten minutes on a skill you can put in the hands of the people who count on you — one that stays theirs, works in any assistant, and outlives whichever app is popular this year.

\$4,800 funds one cohort of 16–20 — your staff, your members, or a mix · they never pay a cent · the price is on this slide because there is no follow-up sequence

01 · WHAT THEY'RE ACTUALLY HOLDING

Hope this email finds you _____

You filled that in before I finished the sentence. **So does the assistant everyone in your building is already using** — for every word it writes. It is autocomplete at enormous scale: it continues with what most people would say.

A simplification an expert accepts: predicting the most likely next word is the literal training objective.

02 · THE OBJECTIVE, PLAINLY

It always bets on **average**.

Not a flaw — the definition. The most likely word is the most common one, so every draft that comes back to the folks you serve **starts life as the most common version of the thought**. They can beat that. Nobody has ever shown them how.

03 · THE HALF NOBODY MENTIONS

And it **reads** them the same way.

Interpretation is prediction too. Handed a request from the coordinator who runs your Tuesday programs, it hears the most likely meaning — the **average asker's** meaning. What it misreads is exactly what you value in her: her judgment, her shorthand, the fifteen years of things she knows that most people don't. And it never says so. It answers a question nobody asked, confidently.

Rewording alone swings measured accuracy up to 76 points (Sclar et al., ICLR 2024). Models misread — and covertly penalize — non-standard dialects (Hofmann et al., Nature 2024). It does not read everyone equally well — and the people it reads worst are often the ones you are working hardest to help.

04 · THEY'VE SEEN THIS MOVIE

A profile about you, written by **someone else.**

Twenty years of web cookies: a dossier nobody authored, kept for someone else's purposes, invisible to the person it described. **We ended up regulating it.** Everybody remembers how that felt.

Every one of them already has one.

- AI memory is **on by default** on all five major services surveyed.
- **Inferred, not written by them** — and it rewrites itself over time (OpenAI calls the process "dreaming").
- At Microsoft: **exempt from retention policy, not audit-logged, not restrictable** — their own documentation.
- At Perplexity Enterprise: those memories **belong to the organization**, and access ends when the person leaves.

Receipts cited in "Durable Context" (Young, 2026) — myworkbrain.org. Nobody wrote the profile their assistant reads them through — not the person it describes, not you. Most have never laid eyes on it.

06 · SOMETHING YOU CAN GIVE THEM TONIGHT

Three questions. No cost, no sign-up.

"What do you remember about me?"

Count what is flat wrong.

"Describe how I communicate."

See whether you recognize the person it describes.

"What matters most in my work?"

Watch it answer with everyone's answer.

Send them to your staff, post them for your members, or just try them yourself tonight. Whatever happens on your screen is happening on theirs too — every day, quietly. Where each service keeps its memory page: the field guide at the end.

The honest ledger.

What that memory does well

Zero effort — it fills itself. Continuity across chats. The mundane, reliably: tone, recurring facts, updates nobody would think to record. Concede every bit of it; it's true.

What it cannot do

Confirm anything with the person. Tell them when it rewrote itself. Move with the person to another assistant — or another employer. Read what is distinctive about them as anything but noise.

Let them use it — and don't let it be the **only** thing that speaks for them. It is the model's *reading* of someone, stored: each misread shapes the next one. **A memory nobody authored is a lens nobody can focus.**

08 · THE MEASURED GROUP EFFECT

Individually faster. Collectively identical.

The research finding on AI assistance: each person's output improves, and **the group's diversity shrinks**. Ten people produce ten versions of the same thing — and none of them chose that.

Doshi & Hauser, Science Advances 2024 · Padmakumar & He, ICLR 2024 · Shumailov et al., Nature 2024. A measured trajectory, still early — we say so out loud.

09 · WHAT IT COSTS THEM

The day their work stops sounding like **them**.

For someone on your staff: the contribution that used to be recognizably hers stops getting noticed. For anyone you serve who is job-hunting, in a pile where every application is polished by the same tool, **sameness is the rejection**. And the discount lands on your organization's name, never on the software.

10 • ONE QUESTION TO KEEP

The difference has to come from the **people.**

Your competitor's proposal was written by the same model as yours. So was the grant application from the organization across town. **Has anyone given the folks doing that work a way to put themselves back into it?**

11 · THE TWO DEFAULT POLICIES

Ban it

They use it anyway — personal accounts, at home, with nobody to ask.
You lose the oversight and the leverage, and nobody learns a thing.

Let it run

Unsupported adoption walks everyone toward the same house style.
You keep the leverage and quietly spend what made each person distinct.

Both defaults leave everyone to work it out alone. There is a third option, and it is a **teaching** one.

12 · THE THIRD OPTION

Teach them to **author** their own context.

Each person writes a short, plain-text file about their own work — their names for things, their judgment, the standards they hold — and hands it to whatever assistant they already use. **Nothing routes through us. No account. Nothing transmitted.** The tool finally reads and writes from who they actually are.

A skill that belongs to them.

Their files

Plain markdown they own — who they are, how they work, what they repeat. On their device, not ours.

A method

One job they do every week, written down well enough that a tool — or a new coworker — can follow it.

Proof

Their own before-and-after, run with their own eyes, on their own work.

A credential

A certificate at session four, and **Work Brain Certified** later for demonstrated use — verifiable, and it goes on a profile.

It works in a tool that doesn't exist yet, at their next job, in a model nobody has released. **Nothing here expires when the program ends.**

Numbers, not testimonials.

A public standard

Four checks, five minutes, no judgment call in it. Any skeptic on your staff can run it — and argue with it.

Per-person evidence

Context typed, exchanges needed, whether the result got used — measured before any teaching, and again at the end.

Honest reporting

At 90 and 180 days, always with the response rate attached. We do not claim employment outcomes.

The standard, and the six claims we invite anyone to attack: myworkbrain.org/good

\$4,800. Four evenings. 16–20 seats in your room.

Facilitation, curriculum, printed materials, loaner devices for anyone without one, certificates, and follow-up reporting. You provide a room with wifi and the folks who should be in it. **They never pay, and the tool stays free forever** — no license, no seat fee, nothing that lapses when we leave.

We do not sell to participants — not during the sessions, not after. General liability carried; facilitators background-checked; documentation on request. Taster sessions and train-your-own-facilitator formats exist; the packet has both.

Take it, or keep the three questions.

If this is something you want for the people you serve: request a cohort at myworkbrain.org/civic and the packet, the agreement and available dates come back within two business days. If it isn't: **keep slide six, run it with your staff or your members, and we are square.** There is no follow-up sequence.

APPENDIX A · PUSH BACK ON US

"The homogenization research is early."

It is, and we say so on the slide. What isn't early: memory on by default, profiles nobody authored, and the before-and-after each person measures on their own work. The case never rests its weight on the scary slide.

"This sounds anti-AI. We're trying to drive adoption."

Four evenings of hands-on building with each person's own assistant. The foundations hour exists so the adoption is informed. We teach use — the caution is the seatbelt, not an argument against the car.

"The folks I serve need results now, not voice lessons."

Everyone's output is polished now, and readers pattern-match the house style — in hiring, in sales, in grant review. When polish is table stakes, what survives is whose judgment shows. That is an output argument, not a taste argument.

APPENDIX B · THE HONEST CONCESSIONS

"Isn't this prompt training with extra steps?"

Prompt phrasing doesn't transfer between assistants — the research is clear. An owned file does. The room demonstrates the difference to itself, in five minutes, on their own work.

"You teach people not to sound like the machine... using the machine."

The same way a writing class uses a word processor. The tool drafts; the person and their file decide. Improve what they do — don't replace it.

What we concede without a fight

Average is fine for boilerplate; the file exists for the work that carries someone's name. And a context file steers voice — it does not guarantee it. The promise is authorship plus evidence, never magic.

Where the memory lives.

Claude

Settings → memory. Viewable by topic; import and export are documented.

ChatGPT

Settings → Personalization → Memory. A summary page; background synthesis keeps updating it.

Gemini

Saved Info. Note that its memory and its Gems are mutually exclusive surfaces.

Microsoft 365 Copilot

Settings → personalization. Stored in the person's Exchange mailbox; on by default.

Perplexity

Settings → memory. In Enterprise, those memories belong to the organization.

Paths current as of September 2026 — vendors move them; search "<service> memory settings" if one has drifted. And remember: the page shows only the part they show you.